

End-to-End Computer Vision Project

Introduction

The goal of this project is to design, implement, and evaluate an **end-to-end computer vision system** addressing a well-defined problem, preferably inspired by a **real-world or industrial application**. Students are expected to apply one or more **deep learning techniques for computer vision** covered during the course.

Each project must incorporate **at least one** of the following modules:

- Semantic segmentation
- Object detection
- Super-resolution and/or image denoising

Projects should go beyond isolated model training and instead demonstrate a **complete pipeline**, including data preparation, model design, training, evaluation, and critical analysis of results.

Evaluation Criteria

The project will be evaluated based on three complementary components:

1. Code (Implementation)

Students are expected to deliver a **fully functional and reproducible codebase**.

The code should:

- Implement the complete end-to-end pipeline
- Be well-structured, modular, and readable
- Use appropriate deep learning frameworks and best practices
- Include clear instructions to reproduce the experiments
- Allow correct execution of training and evaluation scripts

Code quality, clarity, and reproducibility will be explicitly assessed.

2. Written Report

Students must submit a **paper-style technical report** describing the project.

The report should:

- Clearly define the problem and its relevance
- Review related work and position the proposed approach
- Describe the dataset, methodology, and experimental setup
- Present quantitative and qualitative results
- Include a critical discussion of limitations and failure cases
- Follow a clear and coherent structure, using precise technical language

The report should reflect **scientific rigor**, not just implementation details.

3. Oral Presentation

Each project will be presented orally to the class.

The presentation should:

- Clearly explain the problem, methodology, and results
- Justify design choices and evaluation strategies
- Demonstrate understanding of the underlying techniques
- Address questions and critically discuss limitations

Both **technical content** and **clarity of communication** will be considered.

Paper-Style Report Template

1. Title
 - Concise and descriptive
2. Abstract (150-250 words)
 - Problem statement and context
 - Proposed method
 - Experimental setup
 - Key quantitative results
 - Main contributions
3. Introduction
 - Background and motivation
 - Objectives of the project
 - Explicit contributions (bullet points)
4. Related Work
 - Overview of relevant classical and deep learning approaches
 - Positioning of the proposed work with respect to existing methods
5. Problem Formulation
 - Formal definition of the task
 - Input/output specification
 - Assumptions and constraints
 - Evaluation metrics

6. Dataset
 - Dataset description and origin
 - Annotation process
 - Train/validation/test split
 - Preprocessing steps
 - Data augmentation strategies
7. Methodology
 - Overall approach
 - Pipeline diagram (mandatory)
 - Description of each module
 - Justification of design choices
8. Experimental Framework
 - Frameworks and libraries used
 - Model configurations
 - Training procedure
 - Hardware setup
 - Reproducibility details
9. Experimental study
 - Baseline models
 - Quantitative results
 - Ablation studies
 - Qualitative results
10. Conclusions and Future Work
 - Summary of findings
 - Future research directions
11. References
 - Use a consistent citation format